

Lab Assignment NO 4

AIM:-

Create a Linear Regression Model using Python/R to predict home prices using Boston Housing

Dataset (<https://www.kaggle.com/c/boston-housing>). The Boston Housing dataset contains

information about various houses in Boston through different parameters. There are 506 samples

and 14 feature variables in this dataset.

The objective is to predict the value of prices of the house using the given features

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
x=np.array([95,85,80,70,60])
y=np.array([85,95,70,65,70])
```

```
model=np.polyfit(x, y, 1)
```

```
model
```

```
array([ 0.64383562, 26.78082192])
```

```
predict=np.polyld(model)
predict(65)
```

```
68.63013698630137
```

```
y_pred= predict (x)
y_pred
```


```
array([87.94520548, 81.50684932, 78.28767123, 71.84931507,
65.4109589 ])
```

```
from sklearn.metrics import r2_score
r2_score(y, y_pred)
```

```
0.4803218090889326
```

```
y_line= model[1] + model[0]*x
plt.plot(x, y_line, c='r')
plt.scatter(x, y_pred)
plt.scatter(x, y, c='r')
```

```
<matplotlib.collections.PathCollection at 0x1e79c2ba890>
```



```
#import numpy as np
#import pandas as pd
#import matplotlib.pyplot as plt

from sklearn.datasets import fetch_openml
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
housing

{'data': array([[ 8.3252      , 41.        , 6.98412698, ...,
2.55555556,
                37.88      , -122.23      ],
                [ 8.3014      , 21.        , 6.23813708, ...,
2.10984183,
                37.86      , -122.22      ],
                [ 7.2574      , 52.        , 8.28813559, ...,
2.80225989,
                37.85      , -122.24      ],
                ...,
                [ 1.7        , 17.        , 5.20554273, ...,
2.3256351 ,
                39.43      , -121.22      ],
                [ 1.8672      , 18.        , 5.32951289, ...,
2.12320917,
                39.43      , -121.32      ],
                [ 2.3886      , 16.        , 5.25471698, ...,
```

```

2.61698113,
    39.37, -121.24 ]]),
'target': array([4.526, 3.585, 3.521, ..., 0.923, 0.847, 0.894]),
'frame': None,
'target_names': ['MedHouseVal'],
'feature_names': ['MedInc',
    'HouseAge',
    'AveRooms',
    'AveBedrms',
    'Population',
    'AveOccup',
    'Latitude',
    'Longitude'],
'DESCR': '.. _california_housing_dataset:\n\nCalifornia Housing
dataset\n-----\n\n**Data Set Characteristics:**\n
\n    :Number of Instances: 20640\n\n    :Number of Attributes: 8
numeric, predictive attributes and the target\n\n    :Attribute
Information:\n        - MedInc            median income in block group\n
- HouseAge        median house age in block group\n        - AveRooms
average number of rooms per household\n        - AveBedrms average
number of bedrooms per household\n        - Population    block group
population\n        - AveOccup    average number of household
members\n        - Latitude    block group latitude\n        -
Longitude    block group longitude\n\n    :Missing Attribute Values:
None\n\nThis dataset was obtained from the StatLib repository.\n
https://www.dcc.fc.up.pt/~ltorgo/Regression/cal_housing.html\n\nThe
target variable is the median house value for California districts,\n
expressed in hundreds of thousands of dollars ($100,000).\n\nThis
dataset was derived from the 1990 U.S. census, using one row per
census\nblock group. A block group is the smallest geographical unit
for which the U.S.\nCensus Bureau publishes sample data (a block group
typically has a population\nof 600 to 3,000 people).\n\nA household is
a group of people residing within a home. Since the average\nnumber of
rooms and bedrooms in this dataset are provided per household, these\ncolumns may take surprisingly large values for block groups with few
households\nand many empty houses, such as vacation resorts.\n\nIt can
be downloaded/loaded using the\n
func:`sklearn.datasets.fetch_california_housing` function.\n\n..
topic:: References\n\n    - Pace, R. Kelley and Ronald Barry, Sparse
Spatial Autoregressions,\nStatistics and Probability Letters, 33
(1997) 291-297\n'}

data = pd.DataFrame(fetch_california_housing().data)

data.columns = fetch_california_housing().feature_names

data.head()

MedInc HouseAge AveRooms AveBedrms Population AveOccup
Latitude \

```

0	8.3252	41.0	6.984127	1.023810	322.0	2.555556
37.88						
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842
37.86						
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260
37.85						
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945
37.85						
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467
37.85						

	Longitude
0	-122.23
1	-122.22
2	-122.24
3	-122.25
4	-122.25

```
df=pd.DataFrame(housing.data, columns=housing.feature_names)
```

```
df
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup
Latitude \						
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556
37.88						
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842
37.86						
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260
37.85						
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945
37.85						
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467
37.85						
...
...						
20635	1.5603	25.0	5.045455	1.133333	845.0	2.560606
39.48						
20636	2.5568	18.0	6.114035	1.315789	356.0	3.122807
39.49						
20637	1.7000	17.0	5.205543	1.120092	1007.0	2.325635
39.43						
20638	1.8672	18.0	5.329513	1.171920	741.0	2.123209
39.43						
20639	2.3886	16.0	5.254717	1.162264	1387.0	2.616981
39.37						

	Longitude
0	-122.23
1	-122.22

```
2      -122.24
3      -122.25
4      -122.25
...      ...
20635   -121.09
20636   -121.21
20637   -121.22
20638   -121.32
20639   -121.24
```

```
[20640 rows x 8 columns]
```

```
data['PRICE'] = housing.target
```

```
data.isnull().sum()
```

```
MedInc      0
HouseAge    0
AveRooms    0
AveBedrms   0
Population  0
AveOccup    0
Latitude    0
Longitude   0
PRICE       0
```

```
dtype: int64
```

```
x = data.drop(['PRICE'], axis = 1)
```

```
y = data['PRICE']
```

```
from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size
=0.2, random_state = 0)
```

```
import sklearn
```

```
from sklearn.linear_model import LinearRegression
```

```
lm = LinearRegression()
```

```
model=lm.fit(xtrain, ytrain)
```

```
ytrain_pred = lm.predict(xtrain)
```

```
ytest_pred = lm.predict(xtest)
```

```
df=pd.DataFrame(ytrain_pred,ytrain)
```

```
df=pd.DataFrame(ytest_pred,ytest)
```

```
from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(ytest, ytest_pred)
print(mse)


0.5289841670367192

mse = mean_squared_error(ytrain_pred,ytrain)
print(mse)

0.5234413607125448


plt.scatter(ytrain ,ytrain_pred,c='blue',marker='o',label='Training
data')

<matplotlib.collections.PathCollection at 0x1e79e542c90>
```




```
plt.scatter(ytest,ytest_pred ,c='lightgreen',marker='s',label='Test
data')

<matplotlib.collections.PathCollection at 0x1e79e5387d0>
```



```
plt.scatter(ytrain ,ytrain_pred,c='blue',marker='o',label='Training
data')
plt.scatter(ytest,ytest_pred ,c='lightgreen',marker='s',label='Test
data')
plt.xlabel('True values')
plt.ylabel('Predicted')
plt.title("True value vs Predicted value")
plt.legend(loc= 'upper left')
plt.plot()
plt.show()
```



Srushti Mhaske B2_13231